

## Beton Project – Section 4: Technical architecture & ML thinking

This PDF corresponds to Section 4 of the expanded markdown in ChatGPT. It outlines Beton's API■first and on■prem■friendly design philosophy, where every core capability is exposed through stable APIs and can be deployed inside a customer's own infrastructure or VPC. The section emphasizes that Beton is deliberately built to satisfy security and compliance needs, not as an afterthought but as a first■class constraint.

On the ML side, the document explains that Beton is moving towards a modular agent architecture. Core LLMs perform reasoning while a tool/MCP layer mediates access to CRM, documents, email, and enrichment flows. Specialized agents focus on research, note■taking, deal strategy, and workflow execution, with a higher■level orchestrator agent coordinating them through a meta■chat interface.

Section 4 also introduces the idea of reinforcement learning over sales outcomes, where agents can be fine■tuned based on which behaviors correlate with faster cycles and higher win rates. It briefly touches on federated RL as a future direction, allowing learning to be shared across customers without compromising data privacy. Altogether, this section gives a high■level but concrete picture of how Beton's AI and infra will be structured.