

Beton Project – Section 9: Broader conceptual work around agents & infra

This PDF corresponds to Section 9 of the expanded markdown in ChatGPT. It focuses on your conceptual work around federated reinforcement learning and the high-level on-prem AI sales stack design. The section explains why federated RL is appealing for Beton: it allows learning from behavior across customers without centralizing raw CRM data in one place.

It outlines how local agent policies could be trained per customer, using rewards based on deal velocity, win rate, and other sales outcomes. Those local updates could then be aggregated into a global model that improves the default behavior of research and workflow agents for everyone. In doing so, Beton can gradually become a “collective intelligence” for B2B sales patterns while still respecting privacy and compliance boundaries.

The section also recounts the high-level architecture you turned into a YC-style pitch: local or VPC-hosted models, a strong orchestration layer, CRM and enrichment infra, and an agent layer that sits on top. This architecture is meant to be explainable to non-engineers yet robust enough for serious technical buyers. In that sense, Section 9 is as much about narrative clarity as it is about technical design.